
Speech Recognition using Spiking Neural Networks on GPU

Swapnil Ahlawat

ID No. 2018A7PS0178G

Department of Computer Science and Information Systems
f20180178@goa.bits-pilani.ac.in

Neelay Shah

ID No. 2018A8PS0400G

Department of Electrical and Electronics Engineering
f20180400@goa.bits-pilani.ac.in

Under the supervision of

Dr. Basabdatta Sen Bhattacharya

Department of Computer Science and Information Systems

November 26, 2020

Abstract

Presently, artificial neural networks(ANNs) are the mainstream modelling technique for automatic speech recognition [1]. A conventional ANN features a deep multi-layer architecture which makes use of multiple matrix computations for processing data. The rapid progress in the integration of voice interfaces has been viable on account of the remarkable performance of the ASR systems using ANNs for acoustic modeling

The performance gains however, come with immense computational requirements often due to the time-synchronous processing of input audio signals. These models are computationally intensive and memory inefficient to operate as compared to the biological brains. Alternatively, event-driven models such as spiking neural networks (SNNs) inspired by the human brain have attracted ever-growing attention in recent years. Unlike ANNs, asynchronous and event-driven information processing of SNNs resembles the computing paradigm that observed in the human brains, whereby the energy consumption matches the activity levels of sensory stimuli.

The aim of this project is to implement and simulate biologically plausible SNNs on GPU compute for a speech recognition task. We attempt to train an SNN capable of classifying speech signals into one of 10 digits (0-9). We use a supervised Spike-timing-dependent plasticity (STDP) learning rule to train network to differentiate between speech signals belonging to different categories of digits.

Acknowledgement

We would like to express our sincere gratitude to Professor Basabdatta Sen Bhattacharya for her constant guidance and support during the course of the project.

We are grateful to Mr. Sounak Dey from TCS Research who has overseen our work during this project and has provided his invaluable inputs during the process.

We would also like to thank Aiswarya Subramanian, an alumni of BITS Pilani Goa for her help and guidance.

We extend our gratitude to Dr. James Knight from the University of Sussex, co-creator of PyGeNN (a Python framework for neuronal network simulation), for his valuable advice on how to use the framework.

We are also grateful to Mr. Shreenivas A Naik, member of technical staff at BITS Pilani Goa, who helped us in setting up and accessing GPU on a remote server.

Contents

<i>Abstract</i>	i
<i>Acknowledgement</i>	ii
1 Introduction	1
2 Izhikevich neuron model	2
3 Spike Time Dependent Plasticity	3
3.1 Hebbian and Anti-hebbian STDP	3
4 PyGeNN - A neuronal network simulation framework	4
5 Current Work	6
5.1 Observing behaviour of Izhikevich neurons with custom parameters . . .	6
5.2 Feature extraction from speech dataset	7
5.3 Network Architecture	7
5.4 Feeding data into input neurons	8
5.5 Implementing STDP in PyGeNN	8
5.6 Teacher supervision in PyGeNN	9
5.7 Network training	10
6 Conclusion and Future Work	12
<i>Appendix</i>	13
<i>References</i>	21

1 Introduction

Spiking neural networks are computational models that simulate neural behavior in biological systems. Like biological networks, they are driven by discrete spike trains whose instantaneous frequency encodes data. These spikes, by either polarisation or depolarisation, will alter the membrane potential of neurons. Dynamics of the neuron are decided by a set of differential equations that relate membrane voltage with conductance and synaptic current. If a neuron's membrane potential exceeds a certain threshold value, it spikes, passing the integrated information to the next neuron. The influence of a synapse on the membrane voltage of the next neuron can be of two types: excitatory and inhibitory. Excitatory synapses depolarise the neuron, pushing it towards the threshold voltage while inhibitory synapses polarise, thereby reducing the likelihood of a spike. Inhibitory neurons are generally useful to regulate spiking frequency and increase more information contained per spike of the postsynaptic neuron.

There are various neuron models that mimic fairly accurately, the dynamics of biological neurons. The Hodgkin-Huxley model is a conductance-based model that uses a set of non-linear differential equations to explain the electrical characteristics of neurons. While it is a powerful model which is biophysically meaningful, it is very complex and computationally expensive to simulate. Hence, simple spiking neuron models are favoured for applications. A simple model which is widely used is Integrate-and-Fire model. Integrate-and-Fire model considers the neuron like an R-C circuit with a current $I(t)$ flowing through the circuit. This current can be simulated either by external current or presynaptic input spikes. option recently has been the Izhikevich neuron model since it is both accurate and computationally efficient for simulation purposes We focus on the Izhikevich neuron model in the next section.

2 Izhikevich neuron model

The Izhikevich's Neuron model is a simple neuron model which combines the biological plausibility, accuracy, scalability and flexibility of the Hodgkin-Huxley model and the computational inexpensiveness of the Integrate and Fire Neuron model.

$$\frac{dv(t)}{dt} = 0.04v^2(t) + 5v(t) + 140u(t) + Ip_{sc}(t) + Idc$$

$$\frac{du(t)}{dt} = a(bv(t) - u(t))$$

If $v(t) > V_{peak}$, then $v(t) \leftarrow c$; $u(t) \leftarrow u(t) + d$

where a , b , c , d are parameters that affect the dynamics of the model and can be chosen to obtain different kinds of spiking behavior. $u(t)$ is the recovery variable $v(t)$ is the membrane voltage. Ip_{sc} is the post-synaptic current and Idc is the DC bias current stimulus

3 Spike Time Dependent Plasticity

Spike timing dependent plasticity (STDP) is a biological process that adjusts the strength of connections between neurons in a network. The process adjusts the connection strengths based on the relative timing of a particular neuron's output and input action potentials (or spikes). An increase in the strength of the connection is called as Long-Term Potentiation(LTP) whereas a decrease is called Long-Term Depression(LTD)

3.1 Hebbian and Anti-hebbian STDP

In the case of Hebbian STDP, if the postsynaptic spike is generated immediately after receiving the presynaptic spike, the presynaptic spike has a causal role in the output neuron firing. The synaptic weight is thus increased (LTP). Conversely, if a postsynaptic spike occurs before the presynaptic spike, the strength is reduced (LTD), as seen in the equation below

$$\Delta w_{ji} = \begin{cases} Ae^{\frac{(-|t_j-t_i|)}{\tau+}}, & t_j - t_i > 0, A > 0. \\ Be^{\frac{(-|t_j-t_i|)}{\tau-}}, & t_j - t_i < 0, B < 0. \end{cases} \quad (1)$$

In the above, the first case (Case 1) covers LTP and the second case (Case 2) covers LTD. Both cases are decaying exponentials that decay with the distance between and pre and postsynaptic spikes. $A > 0$ and $B < 0$ scale the amplitude of the exponential, and $\tau+$ and τ are the respective time constants

Anti-hebbian STDP is simply the reverse of Hebbian STDP

$$\Delta w_{ji} = \begin{cases} Be^{\frac{(-|t_j-t_i|)}{\tau-}}, & t_j - t_i > 0, B < 0. \\ Ae^{\frac{(-|t_j-t_i|)}{\tau+}}, & t_j - t_i < 0, A > 0.. \end{cases} \quad (2)$$

4 PyGeNN - A neuronal network simulation framework

GPU-enhanced Neuronal Networks (GeNN) is a software package to enable neuronal network simulations on NVIDIA GPUs by code generation. It is a cross-platform C++ library for generating optimized CUDA code for GPU accelerated spiking neural network simulations.

Python interface to GeNN (PyGeNN) is an abstraction on top of GeNN that allows users to access all GeNN features from Python

Given below is an example code snippet highlighting how an Izhikevich neuron can be simulated in 4 regimes for 100 ms using PyGeNN -

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from pygenn.genn_model import GeNNModel
4
5 # Create a single-precision GeNN model
6 model = GeNNModel("float", "pygenn")
7
8 # Set simulation timestep to 0.1ms
9 model.dT = 0.1
10
11 # Initialise IzhikevichVariable parameters
12 izk_init = {"V": -65.0,
13             "U": -20.0,
14             "a": [0.02, 0.1, 0.02, 0.02],
15             "b": [0.2, 0.2, 0.2, 0.2],
16             "c": [-65.0, -65.0, -50.0, -55.0],
17             "d": [8.0, 2.0, 2.0, 4.0]}
18
19 # Add neuron populations and current source to model
20 pop = model.add_neuron_population("Neurons", 4, "IzhikevichVariable",
21                                  {}, izk_init)
22 model.add_current_source("CurrentSource", "DC", "Neurons",
23                           {"amp": 10.0}, {})
24
25 # Build and load model
26 model.build()
27 model.load()
28
29 # Create a numpy view to efficiently access the membrane
30 voltage from Python
31 voltage_view = pop.vars["V"].view
32
33 # Simulate
34 v = None
35 while model.t < 100.0:
36     model.step_time()
37     model.pull_state_from_device("Neurons")
38     v = np.copy(voltage_view) if v is None
39     else np.vstack((v, voltage_view))
```




```

40
41 # Create plot
42 figure, axes = plt.subplots(4, sharex=True)
43
44 # Plot voltages
45 for i, t in enumerate(["RS", "FS", "CH", "IB"]):
46     axes[i].set_title(t)
47     axes[i].set_ylabel("V [mV]")
48     axes[i].plot(np.arange(0.0, 200.0, 0.1), v[:,i])
49 axes[-1].set_xlabel("Time [ms]")

```

Figure 1: Membrane potential vs Simulation Time



5 Current Work

We aim to replicate the results of *A spiking network that learns to extract spike signatures from speech signals* [2]. The method described in the paper uses a compact, non-recurrent SNN consisting of Izhikevich neurons equipped with STDP for a spoken digit recognition task.

We are making use of the Free Spoken Digit Dataset (FSDD) - an open speech dataset consisting of recordings of spoken digits in English by 6 native English speakers.


Described hereon are the different steps we have completed towards trying to replicate the previous speech recognition work -

5.1 Observing behaviour of Izhikevich neurons with custom parameters

PyGeNN allows users the flexibility of simulating Izhikevich neurons with custom parameters. We use the values of parameters mentioned in [2] in order to have a faithful comparison when trying to replicate results.

We simulated a single Izhikevich neuron for a range of current values and for a time duration of 1000 ms with a timestep of 1 ms to determine the current sweep for which the neurons display desired spiking behaviour (Figure 2). We observed the frequency of spikes for the neuron for the range of current values. Simultaneously, we also recorded the membrane potential of the neuron to keep track of the model dynamics.


Figure 2: Spiking frequency vs Current



We observed that the neuron starts spiking at an input current of 52 nA, and the spiking frequency saturates at a current value of 59000 nA. We checked the membrane potential over time for the neuron to confirm whether 59000 nA was a suitable upper

bound of input current. However, We noticed irregularities in the action potentials of the neuron for a current value of greater than 52000 nA. Hence, we decided to use 52000 nA as the upper bound of input current to be fed to the Izhikevich neurons during the speech recognition task, with 52 nA being the lower bound.

Figure 3: Membrane potential vs Simulation time



5.2 Feature extraction from speech dataset

We follow the same procedure as [2] for extracting representative feature vectors to be used as input to the SNN from raw speech signals. The speech signal is divided into small overlapping time sections called speech frames. A fixed number of frames, N , is used for each spoken digit. The length, L , of a spoken digit varies from 500 to 1000 ms, and was divided into $N=40$ frames with 50% overlap to support a frame length of 10-50 ms. The 50% overlap captures the temporal characteristics of the changing spectrum of the speech signal.

After framing, a small feature vector based on the frame's frequency spectrum is extracted. The spectrum values are calculated for all the frames temporally to represent the speech signal spectrogram. A frame encompassing an R Hz frequency range can be divided into M frequency bands. We use $R = 4000$ Hz and $M = 5$ to produce a feature vector of length $N*M = 200$ for each spoken digit. The number of filter banks ($M = 5$) is small enough to create a minimal SNN.

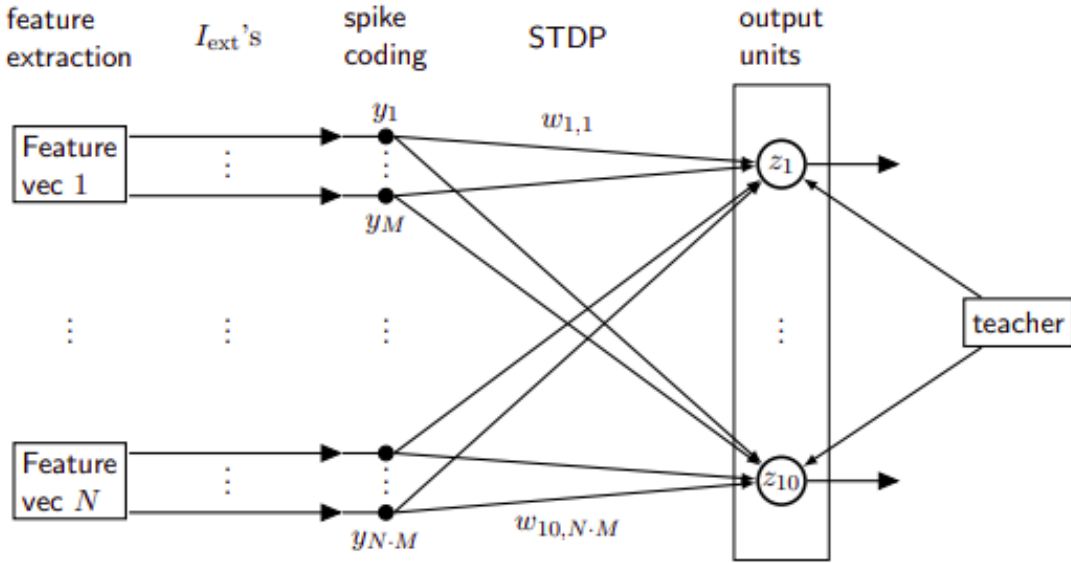
5.3 Network Architecture

We use Izhikevich neurons in our spiking network. The speech signal is divided into 40 input vectors and each input vector has 5 features. Each of these are connected to one neuron. Therefore, the input layer has 200 neurons (denoted by y unit layer).

The output layer has 10 neurons (denoted by z unit layer). Each unit corresponds to one of the ten spoken digit categories (class labels). The y units are fully connected to the z units, i.e., there is no hidden layer.

Finally, there is a supervision signal (called a teacher) that monitors the z units in order to determine the form of the STDP used in training. The teacher determines which z units undergo Hebbian versus anti-Hebbian STDP. During training, whenever a z unit emits a spike, it undergoes some form of STDP. If the z unit represents the target category, then it undergoes Hebbian STDP. Otherwise, it undergoes anti-Hebbian STDP. The teaching signal is only used for the training phase.

Figure 4: Network Architecture



5.4 Feeding data into input neurons


We employ an input layer of 200 Izhikevich neurons in our network; one neuron for each element in the feature vector of a speech signal. The values of the features are used as constant DC current sources for the network simulation. We normalize the feature values between the lower and upper bounds determined previously.

We supply one sample speech signal as input and simulate the network for a period of 100 ms. We observe the behaviour of the 200 input neurons in the form of a raster plot.(Figure 5)

5.5 Implementing STDP in PyGeNN

We implemented both hebbian and anti-hebbian forms of STDP in PyGeNN for a spiking neural network of 2 neuron layers. We use the utility of defining custom learning rules

Figure 5: Raster plot of the input (y) layer neurons



PyGeNN provides to implement both the STDP forms. We initialize the synaptic weights between neurons uniformly between 0.1 and 1.0 and limit them to the same upper and lower bounds during the training process as well. We employ the same STDP parameters (in equations (1) and (2)) as [2]. We track the synaptic weights, spike times and membrane potentials of all the neurons in the network using state variables.

5.6 Teacher supervision in PyGeNN

To facilitate the supervision signal, we make use of an 'index' parameter in the network. Every z unit (output neuron) has the aforementioned 'index' parameter which corresponds to what digit that particular z unit represents. For example, the first z unit would have an index of 0 because it corresponds to the digit 0. We also employ a variable called 'label' which denotes which digit category a speech signal belongs to. The 'label' variable is updated every time a new speech signal is being processed.

Every time a z unit spikes, its index is compared to the label of the speech signal currently being processed. If the index matches the label, the 'correct' neuron for that particular signal has spiked and it undergoes hebbian SDTP. If the index of a neuron doesn't match the label, it undergoes anti-hebbian STDP.

5.7 Network training

We simulated the network for 1 speech sample using teacher supervision as described above. We plotted and observed the behaviour of the neurons in the output (z) layer. Figure 6 shows the membrane potential over time of one z unit. Regular spiking behaviour is observed.

Figure 6: Membrane potential of neuron 1 in the output layer (1st z unit)




Figure 7: Raster plot of the output (z) layer neurons

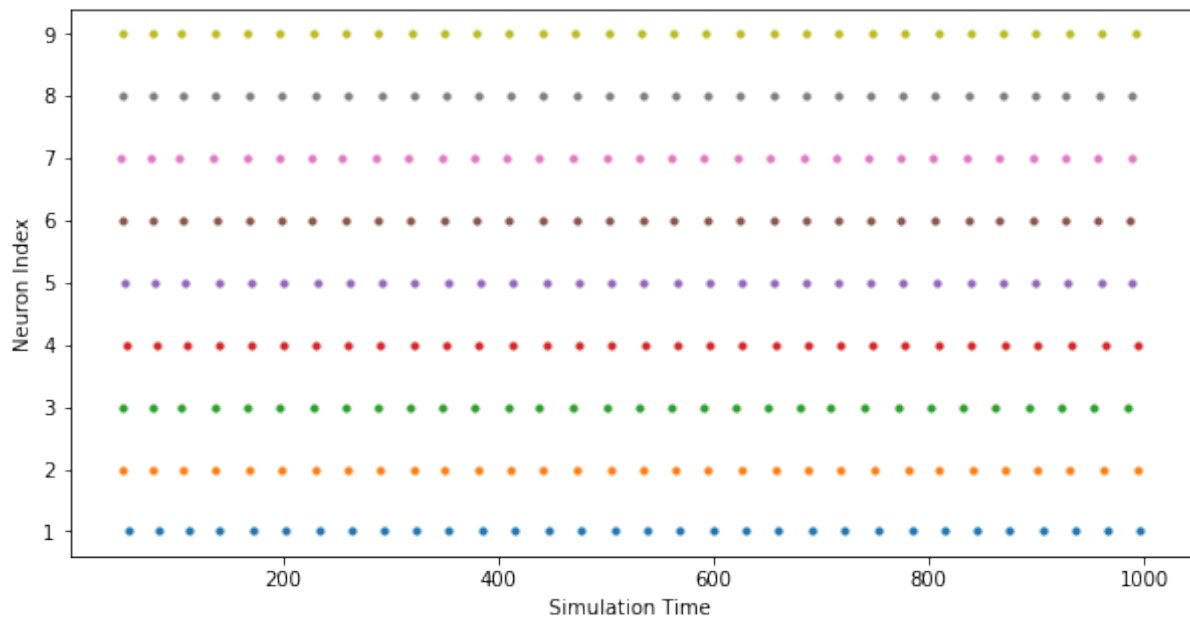


Figure 7 displays a raster plot of the output layer neurons. Since the network hasn't been trained yet with enough training samples to distinguish between speech signals belonging to different categories, all the output layer neurons display similar spiking behaviour. After training, we expect the spiking behaviour of the neurons to be representative of the category of the speech signal being processed.

6 Conclusion and Future Work

During the course of the project, we were able to learn how to simulate neuronal networks on GPU compute using PyGeNN. We understood how spiking neurons process information, how learning takes place in spiking neural networks and what are the ways by which we can decode what an SNN has learnt.

We were able to develop a neuronal network for speech recognition in PyGeNN with an aim of replicating the results in [2]. We tested the working of the network we developed on 1 training sample from a spoken digit speech dataset.

We plan to carry out the following steps here onwards -

- Train the network for the entire training portion of the speech dataset.
- Obtain spike signatures for a test portion of speech signals in the dataset and calculate evaluation metrics to judge performance.

Appendix

PyGeNN code for training the network using supervised STDP -

```
1 # Imports
2
3 import pandas as pd
4 import numpy as np
5 from pygenn import genn_model, genn_wrapper
6 from pygenn.genn_model import (
7     create_custom_neuron_class,
8     create_custom_current_source_class,
9     create_custom_weight_update_class,
10     GeNNModel,
11 )
12 import argparse
13 from os.path import exists, join
14 import os
15
16 # Command line arguments
17
18 parser = argparse.ArgumentParser("Script to train model the SNN using supervised STDP")
19
20 parser.add_argument(
21     "--datafile",
22     type=str,
23     required=True,
24     help="Path to the .npy file containing the speech data",
25 )
26
27 parser.add_argument(
28     "--outdir",
29     type=str,
30     default="output",
31     help="Name of folder where all the output files (membrane potentials etc) should be stored. Folder doesn't need to exist beforehand. (default = output)",
32 )
33
34 parser.add_argument(
35     "--n_samples",
36     type=int,
37     default=1,
38     help="Number of samples in the dataset for which the network should be simulated. (default=1)",
39 )
40
41 args = parser.parse_args()
42
43 # Define custom neuron class
44
45 izk_neuron = create_custom_neuron_class(
46     "izk_neuron",
47     param_names=["a", "b", "c", "d", "C", "k"],
48     var_name_types=[("V", "scalar"), ("U", "scalar")],
```

```

48     sim_code="""
49         $(V)+= (0.5/$(C))*($(k)*$(V) + 60.0)*$(V) + 40.0)-$(U)+$(Isyn
50     )); //at two times for numerical stability
51         $(V)+= (0.5/$(C))*($(k)*$(V) + 60.0)*$(V) + 40.0)-$(U)+$(
52     Isyn));
53         $(U)+=$(a)*$(b)*$(V) + 60.0)-$(U))*DT;
54     """
55     reset_code="""
56         $(V)=$(c);
57         $(U)+=$(d);
58     """
59     threshold_condition_code="$(V) > 30",
60 )
61 # Set the parameters and initial values of variables for the IZK neuron
62     (whose class has been defined above)
63 izk_params = {"a": 0.03, "b": -2.0, "c": -50.0, "d": 100.0, "k": 0.7, "
64     C": 100.0}
65 izk_var_init = {
66     "V": -60.0,
67     "U": 0.0,
68 }
69
70
71 # Define the weight update (supervised STDP learning) rule
72
73 # An additional called parameter called 'index' is created. Each neuron
74     in the ouput layer corresponds to a digit from 0-9. This digit a
75     neuron in the output layer corresponds to is stored in 'index'.
76     Therefore, each post neuron has an index
77
78 # An additional variable called 'label' is created. 'label' is the
79     category (digit:0-9) that the speech signal currently processed
80     belongs to. 'label' is updated every time a new speech signal is
81     being processed
82
83 # 'sim_code' is called when a pre-synaptic spike occurs
84 # 'learn_post_code' is called whena post-synaptic spike occurs
85 # Whether the correct post neuron corresponding to a particular speech
86     signal being processed has spiked is determined by comparing the '
87     label' of the speech signal to the 'index' of the post neruon
88
89 # Logic in 'sim_code': when a pre-synaptic spike occurs -
90 #
91     if
92         the the correct post neuron for the speech signal currently being
93         processed fires , the synaptic strength is reduced (Hebbian
94         learning case 2)
95 #
96     if
97         the the incorrect post neuron for the speech signal currently being

```

```

processed fires , the synaptic strength is reduced (Anti-hebbian
learning case 2)
86
87 # Logic in 'learn_post_code': when a post-synaptic spike occurs -
88 # if
the the correct post neuron for the speech signal currently being
processed fires , the synaptic strength is increased (Hebbian
learning case 1)
89 # if
the the incorrect post neuron for the speech signal currently being
processed fires , the synaptic strength is increased (Anti-hebbian
learning case 1)
90
91 supervised_stdp = create_custom_weight_update_class(
92     "supervised_stdp",
93     param_names=["tauMinus", "tauPlus", "A", "B", "gMax", "gMin", "
index"],
94     var_name_types=[("g", "scalar"), ("label", "scalar")],
95     sim_code="""
96         $(addToInSyn, $(g));
97         scalar dt = $(t) - $(sT_post);
98         if(dt > 0) {
99             if($(index) == $(label)){
100                 scalar timing = exp(-dt / $(tauMinus));
101                 scalar newWeight = $(g) - ($(B) * timing);
102                 $(g) = fmax($(gMin), newWeight);
103             }
104
105             else{
106                 scalar timing = exp(-dt / $(tauPlus));
107                 scalar newWeight = $(g) + ($(A) * timing);
108                 $(g) = fmin($(gMax), newWeight);
109             }
110         }
111
112         """,
113     learn_post_code="""
114         scalar dt = $(t) - $(sT_pre);
115         if (dt > 0) {
116             if($(index) == $(label)){
117                 scalar timing = exp(-dt / $(tauPlus));
118                 scalar newWeight = $(g) + ($(A) * timing);
119                 $(g) = fmin($(gMax), newWeight);
120             }
121
122             else{
123                 scalar timing = exp(-dt / $(tauMinus));
124                 scalar newWeight = $(g) - ($(B) * timing);
125                 $(g) = fmax($(gMin), newWeight);
126             }
127         }
128         """,
129     is_pre_spike_time_required=True,
130     is_post_spike_time_required=True,
131 )

```

```

132
133
134 # Set the initial values of variables of the weight update rule
135
136 # Initialize weights uniformly between 0 and 1
137
138 stdp_var_init = {
139     "g": genn_model.init_var("Uniform", {"min": 0.1, "max": 1.0}),
140     "label": 0.0,
141 } # Initialize label to 0
142
143
144 # Define GeNN model
145
146 model = genn_model.GeNNModel("float", "speech_recognition")
147
148
149 # Add neuron populations
150
151 num_inp_neurons = 200
152 num_output_neurons = 10
153
154 inp_layer = model.add_neuron_population(
155     "input_layer", num_inp_neurons, izk_neuron, izk_params,
156     izk_var_init
157 )
158
159 neuron_layers = [inp_layer]
160
161 for i in range(10):
162     neuron_layers.append(
163         model.add_neuron_population(
164             "output_neuron_" + str(i), 1, izk_neuron, izk_params,
165             izk_var_init
166         )
167     )
168
169 # Create synaptic connections
170
171 # Each synapse group contains synapse connections from all 200 input
172 # neurons to 1 particular neuron in the output layer. 10 such synapse
173 # groups are created, one for every neuron in the output layer
174
175 # Every synapse group contains (belonging to specific output neuron)
176 # contains the 'index' of that neuron
177
178 syn_io = []
179 for i in range(num_output_neurons):
180     syn_io.append(
181         model.add_synapse_population(
182             "synapse_input_output_" + str(i),
183             "DENSE_INDIVIDUALG",
184             genn_wrapper.NO_DELAY,
185             inp_layer,

```

```

182         neuron_layers[i + 1],
183         supervised_stdp,
184         {
185             "tauMinus": 20.0,
186             "tauPlus": 20.0,
187             "A": 0.1,
188             "B": 0.1,
189             "gMax": 1.0,
190             "gMin": 0.1,
191             "index": float(i),
192         },
193         stdp_var_init,
194         {},
195         {},
196         "DeltaCurr",
197         {},
198         {},
199     )
200 )
201
202
203 # Define current source
204
205 current_source = create_custom_current_source_class(
206     "current_source",
207     var_name_types=[("magnitude", "scalar")],
208     injection_code="$ (injectCurrent, $(magnitude));",
209 )
210
211
212 # Create current input
213
214 current_input = model.add_current_source(
215     "input_current", current_source, inp_layer, {}, {"magnitude": 0.0}
216 )
217
218
219 # Set simulation parameters
220
221 timesteps_per_sample = (
222     1000.0 # No. of timesteps one speech signal in the dataset is
223           # presented for
224 )
225 resolution = 1
226
227 # Build and load model
228
229 model.dT = resolution
230 model.build()
231 model.load()
232
233
234 # Load data
235

```

```

236 dataset = np.load(args.datafile, allow_pickle=True).item()
237 data = dataset["data"]
238 labels = dataset["labels"]
239
240
241 # Initialize data structures for variables and parameters we want to
    track
242
243 layer_spikes = [(np.empty(0), np.empty(0)) for _ in enumerate(
    neuron_layers)]
244 layer_voltages = [l.vars["V"].view for l in neuron_layers]
245 current_input_magnitude = current_input.vars["magnitude"].view[:]
246 neuron_labels = [
247     syn_io[neuron].vars["label"].view for neuron in range(
    num_output_neurons)
248 ]
249
250 input_voltage_view = inp_layer.vars["V"].view[:]
251 input_voltage = None
252 output_voltage = {}
253 output_voltage_view = {}
254
255 for index, output_neuron in enumerate(neuron_layers[1:]):
256     output_voltage[index] = None
257     output_voltage_view[index] = output_neuron.vars["V"].view
258
259 synaptic_weights = {}
260 synaptic_weight_views = {}
261 for index, synapse in enumerate(syn_io):
262     synaptic_weights[index] = None
263     synaptic_weight_views[index] = synapse.vars["g"].view[:]
264
265
266 # Simulate
267
268 num_simulation_samples = args.n_samples
269
270 while model.t < timesteps_per_sample * num_simulation_samples:
271
272     timestep_in_example = model.t % timesteps_per_sample
273     sample = int(model.t // timesteps_per_sample)
274
275     if timestep_in_example == 0: # If a new sample is starting to be
        processed
276
277         label = labels[sample]
278         print(f"Processing sample {sample}: {label}")
279
280         current_input_magnitude[:] = data[
281             sample
282         ] # Update the current input for the new sample
283         model.push_var_to_device("input_current", "magnitude")
284
285         for l, v in zip(neuron_layers, layer_voltages):

```

```

287         v[:] = -65.0 # Manually 'reset' voltage
288         l.push_var_to_device("V")
289
290         for index, synapse in enumerate(syn_io):
291             neuron_labels[index] = float(label)
292             synapse.push_var_to_device("label") # Update the 'label'
for the new sample
293
294     model.step_time() # Simulate a timestep
295
296     # record input neurons membrane potential
297
298     model.pull_state_from_device("input_layer")
299     input_voltage = (
300         np.copy(input_voltage_view)
301         if input_voltage is None
302         else np.vstack((input_voltage, input_voltage_view))
303     )
304
305     # record output neurons membrane potential
306
307     for index, output_neuron in enumerate(neuron_layers[1:]):
308         model.pull_state_from_device("output_neuron_" + str(index))
309         output_voltage[index] = (
310             np.copy(output_voltage_view[index])
311             if output_voltage[index] is None
312             else np.hstack((output_voltage[index], output_voltage_view[
index])))
313         )
314
315     # record synaptic weights
316
317     for synapse_index, synapse in enumerate(syn_io):
318         synapse.get_var_values("g")
319         synaptic_weights[synapse_index] = (
320             np.copy(synaptic_weight_views[synapse_index])
321             if synaptic_weights[synapse_index] is None
322             else np.vstack(
323                 (synaptic_weights[synapse_index], synaptic_weight_views
[synapse_index]))
324         )
325     )
326
327     # record spikes
328
329     for i, l in enumerate(neuron_layers):
330         model.pull_current_spikes_from_device(l.name)
331         spike_times = np.ones_like(l.current_spikes) * model.t
332         layer_spikes[i] = (
333             np.hstack((layer_spikes[i][0], l.current_spikes)),
334             np.hstack((layer_spikes[i][1], spike_times)),
335         )
336
337
338 # Create ouput directory

```

```

339
340 if not exists(args.outdir):
341     os.mkdir(args.outdir)
342
343
344 # Save all output files
345
346 pd.DataFrame(input_voltages).to_csv(
347     join(args.outdir, "input_neurons_membrane_potential.csv"), header=
        None, index=None
348 )
349 np.save(join(args.outdir, "output_neurons_membrane_potential"),
        output_voltages)
350 np.save(join(args.outdir, "synaptic_weights"), synaptic_weights)
351 np.save(join(args.outdir, "spikes_data"), layer_spikes)

```


References

- [1] J. Wu, E. Yilmaz, M. Zhang, H. Li, and K. C. Tan, “Deep spiking neural networks for large vocabulary automatic speech recognition,” *Frontiers in Neuroscience*, vol. 14, p. 199, 2020.
- [2] A. Tavanaei and A. S. Maida, “A spiking network that learns to extract spike signatures from speech signals,” *Neurocomputing*, vol. 240, p. 191–199, May 2017.